

CASE STUDY

Image Classification Using K-Nearest Neighbor (KNN)

Batch – 17

22251A0536 – B.Sirisha

22251A0541 – Tanmayi

22251A0548 – K.Asmitha

22255A0560 – R.Ashwitha

DATASET DESCRIPTION

Context:

The **CIFAR-10 dataset** is a widely used benchmark for image classification tasks. It consists of **60,000 color images**, each sized **32×32 pixels**, categorized into **10 distinct classes**. The dataset is used to evaluate and compare machine learning models on visual object recognition.

Classes:

- **airplane, automobile, bird, cat, deer,**
- **dog, frog, horse, ship, truck**

Dataset Structure

- **Training Set:** 50,000 images
- **Test Set:** 10,000 images
- **Image Shape:** Each image is a **3D array** of size **(32, 32, 3)** representing RGB color values

Attributes Used

1. **Images:** Raw color images used as input data
2. **Labels:** Numeric values (0–9) corresponding to the image classes
3. **Preprocessing:**
 - a. **Flattening or feature extraction** methods (e.g., PCA or raw pixel features)
 - b. Optional **normalization** of pixel values (e.g., scaling between 0–1)
4. **Distance Metric:**
 - a. **Euclidean distance** used to measure similarity between images

Model Training and Evaluation

Algorithm Used: K-Nearest Neighbors (KNN)

- **How it works:** KNN classifies an image based on the **majority label among its k-nearest neighbors** in the feature space
- **k value used:** Typically 3 or 5 (tuned via validation)
- **No model training required** – classification is performed during the prediction phase

Why KNN?

- **Simple and intuitive** to understand and implement
- **No explicit training phase** (lazy learning algorithm)
- Effective for **smaller datasets** or as a baseline model

Evaluation Metrics

- **Accuracy:** Measured as the percentage of correctly classified test images
- **Confusion Matrix:** Shows correct and incorrect classifications across all 10 classes
- **Optional Metrics:** Precision, recall, and F1-score per class (for deeper analysis)

Visualization

- **Sample Images:** Visual representation of images from each class
- **Confusion Matrix:** Heatmap to highlight misclassification patterns
- **Predictions:** Random test samples shown with **true vs. predicted labels**

IMPLEMENTATION

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, accuracy_score, confusion_matrix
from keras.datasets import cifar10
import seaborn as sns

(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train = x_train.reshape((50000, -1)) / 255.0
x_test = x_test.reshape((10000, -1)) / 255.0
y_train = y_train.ravel()
y_test = y_test.ravel()

x_train_small = x_train[:5000]
y_train_small = y_train[:5000]
x_test_small = x_test[:1000]
y_test_small = y_test[:1000]

knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(x_train_small, y_train_small)

y_pred = knn.predict(x_test_small)

print("\n--- Classification Report ---\n")
print(classification_report(y_test_small, y_pred))

accuracy = accuracy_score(y_test_small, y_pred)
print(f"Accuracy: {accuracy:.2f}")
```

OUTPUT:

Downloading data from <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz>
170498071/170498071 ————— 4s 0us/step

--- Classification Report ---

	precision	recall	f1-score	support
0	0.27	0.55	0.36	103
1	0.41	0.16	0.23	89
2	0.17	0.45	0.25	100
3	0.29	0.17	0.22	103
4	0.15	0.33	0.21	90
5	0.17	0.07	0.10	86
6	0.37	0.20	0.26	112
7	0.60	0.09	0.15	102
8	0.49	0.52	0.50	106
9	0.56	0.05	0.08	109
accuracy			0.26	1000
macro avg	0.35	0.26	0.24	1000
weighted avg	0.35	0.26	0.24	1000

Accuracy: 0.26

SAMPLE IMAGE PREDICTIONS

```
classes = ['airplane', 'automobile', 'bird', 'cat', 'deer',  
           'dog', 'frog', 'horse', 'ship', 'truck']
```

```
def show_predictions(num=5):  
    for i in range(num):  
        plt.imshow(x_test_small[i].reshape(32, 32, 3))  
        plt.title(f'Actual: {classes[y_test_small[i]]} | Predicted: {classes[y_pred[i]]}')  
        plt.axis('off')  
        plt.show()
```

```
show_predictions()
```

OUTPUT:


Actual: frog | Predicted: deer




Actual: cat | Predicted: deer



Actual: ship | Predicted: ship



Actual: ship | Predicted: ship




Actual: airplane | Predicted: airplane



CONFUSION MATRIX

```
cm = confusion_matrix(y_test_small, y_pred)
plt.figure(figsize=(10, 8))
sns.heatmap(cm, annot=True, fmt='d', xticklabels=classes, yticklabels=classes,
            cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```


OUTPUT:



RESULTS

Example Output (Actual Run):

Accuracy: 0.26

Precision (avg): 0.35

Recall (avg): 0.26

F1-Score (avg): 0.24

CONCLUSION

KNN is a straightforward approach to image classification that performs reasonably well on small datasets. Its simplicity makes it ideal for foundational learning, though it's computationally expensive for large-scale tasks. For higher accuracy, CNN-based models are recommended.